

# Highly Available & Scalable Web Application on AWS (CloudFormation)

This project demonstrates deploying a highly available, auto-scaling web application on AWS using a CloudFormation YAML template.

## Architecture Overview

- **EC2 Instances** (Amazon Linux 2)
- **Application Load Balancer** (ALB)
- **Auto Scaling Group** (2 to 5 instances)
- **Launch Template** with UserData to install a web server
- **Public Subnets** in 2 Availability Zones
- **Health Checks** based on ALB / path
- **Scaling Policies:**
  - Scale-out at **>60% CPU**
  - Scale-in at **<30% CPU**

## Steps Performed

### 1. Launch CloudFormation Stack

- Used the provided YAML file with:
  - AMIID: Amazon Linux 2 AMI
  - KeyName: real (EC2 Key Pair already created)
  - InstanceType: t2.micro

### 2. Access the Web Application

- After successful deployment:
  - Go to **Outputs** in CloudFormation
  - Copy the LoadBalancerDNS value
  - Paste it in your browser to access: `<h1>Welcome to Auto Scaling EC2 Instance</h1>`

### 3. Simulate CPU Load to Trigger Scaling

- SSH into one of the EC2 instances: `stress --cpu 2 --timeout 300`
- This pushes the average CPU above 60% and triggers scale-out.

#### 4. Allow Scale-In

- Wait for stress to finish (or kill the process)
- Wait 5+ minutes (grace period)
- ASG detects CPU below 30% and scales in
- Note: MinSize is 2, so no scale-in below 2

#### Outputs

- **ALB DNS Name:** Public URL for accessing the application
- **ASG Events:** Auto scaling logs visible in EC2 → Auto Scaling Group

#### Files

- **ASG.yaml:** CloudFormation template to deploy the entire architecture



**Welcome to Auto Scaling EC2 Instance**

**Image: Web Page Served by EC2 Instance**




Image: CPU Utilization Triggering Scale-Out



Image: CPU Utilization Dropping After Load